

# 收获，不止 SQL 优化

## 第八章

且慢，学习索引是如何让 SQL 飞

E-Mail:45240040@qq.com

### 目录

- 1.解答第一题..... 2
  - 1.1 索引高度较低.....2
  - 1.2 索引存储列值.....4
  - 1.3 索引本身有序.....4
- 2.解答第二题..... 6
  - 2.1 组合索引需要考虑单列.....6
  - 2.2 组合索引需要考虑回表.....7
  - 2.3 组合索引需要考虑排序.....8
  - 2.4 组合索引需要考虑顺序.....9
- 3.解答第三题..... 10
  - 3.1 环境准备.....11
  - 3.2 分区表下的 max 聚合函数..... 12
  - 3.3 分区表下的 count 聚合函数..... 12
  - 3.3 分区表下的 sum 聚合函数..... 14
  - 3.3 分区表下的 distinct 去重处理.....15

## 1.解答第一题

◆ 说说老师课堂的索引三大特性是什么，能应用在哪些 SQL 上。

### 1.1 索引高度较低

在 SQL 检索数据(SELECT)的时候,索引的高度的不同对检索的效率有明显的差别,数据库访问索引需要读取的数据块通常是索引的高度+1 个数据块数,也就是说索引的高度越高,访问索引需要读取的数据块数越多,效率越差。

```
SQL> conn scott/tiger
Connected.
SQL> create table t1 as select rownum as id ,rownum+1 as id2, rpad('*',1000,'*')
  2  as contents from dual connect by level<=100;

Table created.

SQL> create table t2 as select rownum as id ,rownum+1 as id2, rpad('*',1000,'*')
  2  as contents from dual connect by level<=10000;

Table created.

SQL> create table t3 as select rownum as id ,rownum+1 as id2, rpad('*',1000,'*')
  2  as contents from dual connect by level<=1000000;

Table created.

SQL> create index idx_id_t1 on t1(id);

Index created.

SQL> create index idx_id_t2 on t2(id);

Index created.

SQL> create index idx_id_t3 on t3(id);

Index created.
```

```
SQL> set lines 120
SQL> col index_name for a15
SQL> select index_name, blevel, leaf_blocks, num_rows, distinct_keys,
  2  clustering_factor from user_ind_statistics
  3  where table_name in( 'T1','T2','T3');
```

INDEX_NAME	BLEVEL	LEAF_BLOCKS	NUM_ROWS	DISTINCT_KEYS	CLUSTERING_FACTOR
IDX_ID_T2	1	21	10000	10000	1429
IDX_ID_T1	0	1	100	100	15
IDX_ID_T3	2	2226	1000000	1000000	142858

```
SQL> set autotrace traceonly stat
SQL> select id from t1 where id=1;
```

#### Statistics

```
-----
0 recursive calls
0 db block gets
2 consistent gets
0 physical reads
0 redo size
405 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> select id from t2 where id=1;
```

#### Statistics

```
-----
0 recursive calls
0 db block gets
3 consistent gets
0 physical reads
0 redo size
405 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> select id from t3 where id=1;
```

#### Statistics

```
-----
0 recursive calls
0 db block gets
4 consistent gets
0 physical reads
0 redo size
405 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

从上面的试验中可以看到，随着表索引高度的增加，其查询索引产生的一致性读也随之增加。

## 1.2 索引存储列值

分析一个索引块我们可以知道索引块不仅存储了 **rowid** 信息，而且还存储了索引列的值，那么当我们查询的值正好是在索引里时或者做一些聚合计算（如 **sum,max,min**）时，就可以利用这个特性。

```
SQL> create table t as select * from dba_objects;
Table created.

SQL> create index indx_t on t(object_id);
Index created.

SQL> set autot traceonly exp
SQL> select count(*) from t where object_id is not null;

Execution Plan
-----
Plan hash value: 2692964945

-----
| Id | Operation          | Name  | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT    |       |      1 |    13 |       7 (0)| 00:00:01 |
|  1 | SORT AGGREGATE      |       |      1 |    13 |           |          |
|*  2 |  INDEX FAST FULL SCAN| INDX_T | 10296 | 130K |       7 (0)| 00:00:01 |
-----
```

## 1.3 索引本身有序

从索引的存储结构上可以看到，索引的存储是有序存放的，扫描索引的时候是从根节点开始，经过颈节点到叶子节点，这个特点下索引的范围查询或等值查询，索引只需要扫描一段范围就可得出结果，因为其本身是有范围的，我们可用利用索引这个特点来降低实际查询的排序操作。



```

SQL> create table t as select * from dba_objects;

Table created.

SQL> exec dbms_stats.gather_table_stats(user,'T');

PL/SQL procedure successfully completed.

SQL> set autot traceonly exp
SQL> set autot traceonly stat
SQL> select * from t where object_id >500 order by object_id;

8992 rows selected.

Statistics
-----
      1 recursive calls
       0 db block gets
     122 consistent gets
       0 physical reads
       0 redo size
  413948 bytes sent via SQL*Net to client
    6973 bytes received via SQL*Net from client
     601 SQL*Net roundtrips to/from client
       1 sorts (memory)
       0 sorts (disk)
     8992 rows processed

```

```

SQL> create index idx_t on t(object_id);

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> select * from t where object_id >500 order by object_id;

8992 rows selected.

Statistics
-----
      1 recursive calls
       0 db block gets
    1334 consistent gets
       0 physical reads
       0 redo size
  887877 bytes sent via SQL*Net to client
    6973 bytes received via SQL*Net from client
     601 SQL*Net roundtrips to/from client
       0 sorts (memory)
       0 sorts (disk)
     8992 rows processed

```

## 2.解答第二题



说说用组合索引需要考虑什么问题。

### 2.1 组合索引需要考虑单列

在创建组合索引的时候需要考虑到常会查询的单列索引，因为创建组合索引会有前导列的概念，在查询中最好应该是应到前导列，这样的查询效率会比较高。

```
SQL> create table t as select * from dba_objects;
Table created.
SQL> create index indx_t_1 on t(object_id,object_type);
Index created.
SQL> set autot traceonly exp
SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);
PL/SQL procedure successfully completed.
SQL> select * from t where object_id=1000;
```

查询的值是前导列cost值

```
Execution Plan
-----
Plan hash value: 1825404486

-----
| Id | Operation                      | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
-----
| 0  | SELECT STATEMENT                |           |      1 |    84 |      3 (0)| 00:00:01 |
| 1  | TABLE ACCESS BY INDEX ROWID    | T         |      1 |    84 |      3 (0)| 00:00:01 |
|* 2  | INDEX RANGE SCAN                | INDX_T_1  |      1 |      |      2 (0)| 00:00:01 |
-----
```

Predicate Information (identified by operation id):

```
-----
2 - access("OBJECT_ID"=1000)
```



```

SQL> drop index indx_t_1;

Index dropped.

SQL> create index indx_t_2 on t(object_type,object_id);

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> select * from t where object_id=1000;

```

查询的列不属于前导列cost

```

Execution Plan
-----
Plan hash value: 1601196873

-----
| Id | Operation          | Name | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT   |      |    1 |    84 |    34 (0)| 00:00:01 |
|*  1 | TABLE ACCESS FULL| T    |    1 |    84 |    34 (0)| 00:00:01 |
-----

Predicate Information (identified by operation id):
-----

   1 - filter("OBJECT_ID"=1000)

```

## 2.2 组合索引需要考虑回表

```

SQL> create table t as select * from dba_objects;

Table created.

SQL> create index idx_t on t(object_id);

```

单索引产生了回表访问

```

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> set autot traceonly exp
SQL> select object_id,object_type from t where object_id >50000;

```

```

Execution Plan
-----
Plan hash value: 1594971208

-----
| Id | Operation          | Name | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT   |      |    1 |    11 |     3 (0)| 00:00:01 |
|  1 | TABLE ACCESS BY INDEX ROWID | T    |    1 |    11 |     3 (0)| 00:00:01 |
|*  2 | INDEX RANGE SCAN   | IDX_T |    1 |      |     2 (0)| 00:00:01 |
-----

```



```
SQL> drop index idx_t;

Index dropped.

SQL> create index idx_t on t(object_id,object_type);

Index created.

SQL> select object_id,object_type from t where object_id >50000;

Execution Plan
-----
Plan hash value: 2296882198

-----
| Id | Operation          | Name | Rows | Bytes | Cost (%CPU)| Time |
-----
|  0 | SELECT STATEMENT   |      |     1 |    11 |        2 (0)| 00:00:01 |
|*  1 |  INDEX RANGE SCAN | IDX_T |     1 |    11 |        2 (0)| 00:00:01 |
-----
```

复合索引下执行计划已经看不到回表访问

## 2.3 组合索引需要考虑排序

```
SQL> create table t as select * from dba_objects;

Table created.

SQL> create index idx_t on t(owner,object_type);

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> set autot traceonly stat
SQL> select /*+ index(t,idx_t)*/ from t
  2  order by owner asc,object_type desc;

9471 rows selected.

Statistics
-----
      1  recursive calls
        0  db block gets
     122  consistent gets
        0  physical reads
        0  redo size
  435285  bytes sent via SQL*Net to client
    7325  bytes received via SQL*Net from client
    633  SQL*Net roundtrips to/from client
      1  sorts (memory)
        0  sorts (disk)
   9471  rows processed
```

产生一个排序

```

SQL> drop index idx_t;

Index dropped.

SQL> create index idx_t on t(owner asc,object_type desc);

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> select /*+ index_desc(t,idx_t)*/ from t
      2  order by owner asc,object_type desc;

9471 rows selected.

Statistics
-----
      1  recursive calls
      0  db block gets
     122  consistent gets
      0  physical reads
      0  redo size
  435285  bytes sent via SQL*Net to client
    7325  bytes received via SQL*Net from client
     633  SQL*Net roundtrips to/from client
      0  sorts (memory)
      0  sorts (disk)
     9471  rows processed

```

## 2.4 组合索引需要考虑顺序

```

SQL> drop table t purge;

Table dropped.

SQL> create table t as select * from dba_objects;

Table created.

SQL> create index idx_t on t(object_id,object_type);

Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);

PL/SQL procedure successfully completed.

SQL> select * from t where object_type='TABLE' and object_id >100;

798 rows selected.

Statistics
-----
      1  recursive calls
      0  db block gets
     188  consistent gets
      0  physical reads
      0  redo size
  79069  bytes sent via SQL*Net to client
    967  bytes received via SQL*Net from client
     55  SQL*Net roundtrips to/from client
      0  sorts (memory)
      0  sorts (disk)
     798  rows processed

```

```
SQL> drop index idx_t;
Index dropped.

SQL> create index idx_t on t(object_type,object_id);
Index created.

SQL> exec dbms_stats.gather_table_stats(user,'T',cascade=>true);
PL/SQL procedure successfully completed.

SQL> select * from t where object_type='TABLE' and object_id >100;
798 rows selected.
```

顺序不同产生的一致性读不同

```
Statistics
-----
      1 recursive calls
      0 db block gets
     175 consistent gets
      0 physical reads
      0 redo size
  37495 bytes sent via SQL*Net to client
   967 bytes received via SQL*Net from client
    55 SQL*Net roundtrips to/from client
      0 sorts (memory)
      0 sorts (disk)
   798 rows processed
```

### 3.解答第三题



分区表中的聚合语句有什么特别之处？

首先我们将在分区表中执行一些聚合语句来总结其特别之处：

### 3.1 环境准备

```
SQL> create table range_part_tab (id number,deal_date date,area_code number,
2  nbr number,contents varchar2(4000))
3  partition by range (deal_date)
4  (partition p_201301 values less than (TO_DATE('2013-02-01', 'YYYY-MM-DD')),
5  partition p_201302 values less than (TO_DATE('2013-03-01', 'YYYY-MM-DD')),
6  partition p_201303 values less than (TO_DATE('2013-04-01', 'YYYY-MM-DD')),
7  partition p_201304 values less than (TO_DATE('2013-05-01', 'YYYY-MM-DD')),
partition p_201305 values less than (TO_DATE('2013-06-01', 'YYYY-MM-DD')),
9  partition p_201306 values less than (TO_DATE('2013-07-01', 'YYYY-MM-DD')),
partition p_201308 values less than (TO_DATE('2013-09-01', 'YYYY-MM-DD')),
partition p_201309 values less than (TO_DATE('2013-10-01', 'YYYY-MM-DD')),
partition p_201310 values less than (TO_DATE('2013-11-01', 'YYYY-MM-DD')),
partition p_201311 values less than (TO_DATE('2013-12-01', 'YYYY-MM-DD')),
partition p_201312 values less than (TO_DATE('2014-01-01', 'YYYY-MM-DD')),
partition p_201401 values less than (TO_DATE('2014-02-01', 'YYYY-MM-DD')),
partition p_201402 values less than (TO_DATE('2014-03-01', 'YYYY-MM-DD')),
17 partition p_max values less than (maxvalue)
18 );
```

Table created.

创建分区表并插入数据

```
SQL> alter table RANGE_PART_TAB modify nbr not null;
```

Table altered.

```
SQL> insert into range_part_tab (id,deal_date,area_code,nbr,contents)
2  select rownum,
3  to_date( to_char(sysdate-365,'J')+TRUNC(DBMS_RANDOM.VALUE(0,365)), 'J'),
4  ceil(dbms_random.value(591,599)),
5  ceil(dbms_random.value(18900000001,18999999999)),
6  rpad('*',400,'*')
7  from dual connect by rownum <= 100000;
```

100000 rows created.

```
SQL> commit;
```

Commit complete.

```
SQL> create index idx_part_id on range_part_tab (id) ;
```

Index created.

```
SQL> create index idx_part_nbr on range_part_tab (nbr) local;
```

Index created.

```
SQL> begin
2  dbms_stats.gather_table_stats(ownname=>'SCOTT',
3  tabname=>'RANGE_PART_TAB',
4  estimate_percent=>10,
5  method_opt=>'for all indexed columns',
6  cascade=>true);
7  end;
8  /
```

PL/SQL procedure successfully completed.



### 3.2 分区表下的 max 聚合函数

```
SQL> set autot traceonly ;
SQL> set lines 200
SQL> select max(nbr) max_nbr from range_part_tab partition(p_201305);
```

Execution Plan

Plan hash value: 1219885076

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	8	2 (0)	00:00:01		
1	SORT AGGREGATE		1	8				
2	PARTITION RANGE SINGLE		8398	67184	2 (0)	00:00:01	5	5
3	INDEX FULL SCAN (MIN/MAX)	IDX_PART_NBR	8398	67184	2 (0)	00:00:01	5	5

Statistics

0	recursive calls
0	db block gets
2	consistent gets
0	physical reads
0	redo size
415	bytes sent via SQL*Net to client
384	bytes received via SQL*Net from client
2	SQL*Net roundtrips to/from client
0	sorts (memory)
0	sorts (disk)
1	rows processed

直接可以利用分区索引，系统知道直接去第5分区查询数据

```
SQL> select max(nbr) max_nbr from range_part_tab
2 where deal_date >= TO_DATE('2013-05-01', 'YYYY-MM-DD')
3 and deal_date < TO_DATE('2013-06-01', 'YYYY-MM-DD');
```

Execution Plan

Plan hash value: 4190023598

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	17	145 (1)	00:00:02		
1	SORT AGGREGATE		1	17				
2	PARTITION RANGE SINGLE		21	357	145 (1)	00:00:02	5	5
3	TABLE ACCESS FULL	RANGE_PART_TAB	21	357	145 (1)	00:00:02	5	5

Statistics

1	recursive calls
0	db block gets
530	consistent gets
0	physical reads
0	redo size
415	bytes sent via SQL*Net to client
384	bytes received via SQL*Net from client
2	SQL*Net roundtrips to/from client
0	sorts (memory)
0	sorts (disk)
1	rows processed

虽然能够得到相同结果集，但是利用不到索引，系统扫描了全表才能确定数据来源于第5分区，导致一致性读和cost值较大

### 3.3 分区表下的 count 聚合函数

Count/sum/distinct 等与 max 相同，这里将不再说明，只演示一遍。

```
SQL> select count(*) max_nbr from range_part_tab partition(p_201305);
```

#### Execution Plan

Plan hash value: 296899510

Id	Operation	Name	Rows	Cost	(%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	8	(0)	00:00:01		
1	SORT AGGREGATE		1					
2	PARTITION RANGE SINGLE		8398	8	(0)	00:00:01	5	5
3	INDEX FAST FULL SCAN	IDX PART_NBR	8398	8	(0)	00:00:01	5	5

#### Statistics

```

421 recursive calls
0 db block gets
151 consistent gets
0 physical reads
0 redo size
411 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
17 sorts (memory)
0 sorts (disk)
1 rows processed

```

```
SQL> select count(*) max_nbr
2 from range_part_tab
3 where deal_date >= TO_DATE('2013-05-01', 'YYYY-MM-DD')
4 and deal_date < TO_DATE('2013-06-01', 'YYYY-MM-DD');
```

#### Execution Plan

Plan hash value: 4190023598

Id	Operation	Name	Rows	Bytes	Cost	(%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	9	145	(1)	00:00:02		
1	SORT AGGREGATE		1	9					
2	PARTITION RANGE SINGLE		21	189	145	(1)	00:00:02	5	5
3	TABLE ACCESS FULL	RANGE_PART_TAB	21	189	145	(1)	00:00:02	5	5

#### Statistics

```

1 recursive calls
0 db block gets
530 consistent gets
0 physical reads
0 redo size
411 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed

```



### 3.3 分区表下的 sum 聚合函数

```
SQL> select sum(nbr) max_nbr from range_part_tab partition(p_201305);
```

Execution Plan

Plan hash value: 296899510

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	8	8 (0)	00:00:01		
1	SORT AGGREGATE		1	8				
2	PARTITION RANGE SINGLE		8398	67184	8 (0)	00:00:01	5	5
3	INDEX FAST FULL SCAN	IDX_PART_NBR	8398	67184	8 (0)	00:00:01	5	5

Statistics

```

1 recursive calls
0 db block gets
28 consistent gets
0 physical reads
0 redo size
417 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed

```

```
SQL> select sum(nbr) max_nbr
2 from range_part_tab
3 where deal_date >= TO_DATE('2013-05-01', 'YYYY-MM-DD')
4 and deal_date < TO_DATE('2013-06-01', 'YYYY-MM-DD');
```

Execution Plan

Plan hash value: 4190023598

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	17	145 (1)	00:00:02		
1	SORT AGGREGATE		1	17				
2	PARTITION RANGE SINGLE		21	357	145 (1)	00:00:02	5	5
3	TABLE ACCESS FULL	RANGE_PART_TAB	21	357	145 (1)	00:00:02	5	5

Statistics

```

1 recursive calls
0 db block gets
530 consistent gets
0 physical reads
0 redo size
417 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed

```

### 3.3 分区表下的 distinct 去重处理

```
SQL> select distinct(nbr) from range_part_tab partition(p_201305);

8398 rows selected.

Execution Plan
-----
Plan hash value: 2418110982

-----
| Id | Operation          | Name          | Rows  | Bytes | TempSpc | Cost (%CPU) | Time      | Pstart | Pstop |
-----
| 0  | SELECT STATEMENT   |               | 8398  | 67184 |          | 42 (5)      | 00:00:01 |        |       |
| 1  | HASH UNIQUE        |               | 8398  | 67184 | 280K    | 42 (5)      | 00:00:01 |        |       |
| 2  | PARTITION RANGE SINGLE |               | 8398  | 67184 |          | 8 (0)       | 00:00:01 | 5      | 5     |
| 3  | INDEX FAST FULL SCAN| IDX_PART_NBR | 8398  | 67184 |          | 8 (0)       | 00:00:01 | 5      | 5     |
-----

Statistics
-----
      1 recursive calls
        0 db block gets
      28 consistent gets
        0 physical reads
        0 redo size
  148650 bytes sent via SQL*Net to client
    6533 bytes received via SQL*Net from client
     561 SQL*Net roundtrips to/from client
        0 sorts (memory)
        0 sorts (disk)
    8398 rows processed
```

```
SQL> select distinct(nbr)
2  from range_part_tab
3  where deal_date >= TO_DATE('2013-05-01', 'YYYY-MM-DD')
4  and deal_date < TO_DATE('2013-06-01', 'YYYY-MM-DD');

8398 rows selected.

Execution Plan
-----
Plan hash value: 4092261255

-----
| Id | Operation          | Name          | Rows  | Bytes | Cost (%CPU) | Time      | Pstart | Pstop |
-----
| 0  | SELECT STATEMENT   |               | 21    | 357   | 146 (2)     | 00:00:02 |        |       |
| 1  | HASH UNIQUE        |               | 21    | 357   | 146 (2)     | 00:00:02 |        |       |
| 2  | PARTITION RANGE SINGLE |               | 21    | 357   | 145 (1)     | 00:00:02 | 5      | 5     |
| 3  | TABLE ACCESS FULL | RANGE_PART_TAB | 21    | 357   | 145 (1)     | 00:00:02 | 5      | 5     |
-----

Statistics
-----
      1 recursive calls
        0 db block gets
     530 consistent gets
        0 physical reads
        0 redo size
  148650 bytes sent via SQL*Net to client
    6533 bytes received via SQL*Net from client
     561 SQL*Net roundtrips to/from client
        0 sorts (memory)
        0 sorts (disk)
    8398 rows processed
```