

收获，不止 SQL 优化

第十章

且慢，其他索引应用让 SQL 飞

E-Mail:45240040@qq.com

目录

- 1.说说位图索引的优势和不适合使用的场景.....2
 - 1.1 位图索引的优势.....2
 - 1.1.1 高效即席查询.....2
 - 1.1.2 快速统计条数.....4
 - 1.2 不适合的场景.....6
 - 1.2.1 列修改容易造成锁表.....6
 - 1.2.2 列值重复度低效率低.....7
- 2.简要说说函数索引的原理及好处.....9
 - 2.1 函数索引的原理.....9
 - 2.2 使用函数索引的好处.....9
 - 2.2.1 对部分记录建立函数索引.....9
 - 2.2.1 减少递归调用.....12
- 3.说说使用反向索引、全文索引的应用场景.....13
 - 3.1 反向索引的应用场景.....13
 - 3.2 全文索引的应用场景.....14
 - 3.2.1 全文索引环境准备.....14
 - 3.2.2 创建语法分析器并设置属性.....15
 - 3.2.3 创建表及为表创建全文索引.....15
 - 3.2.4 比较全文索引的性能.....16

1.说说位图索引的优势和不适合使用的场景

1.1 位图索引的优势

位图索引与其他索引不同，它不是存储的索引列的列值，而是以比特位 0、1 的形式存储，所以在空间上它占的空间比较小，相应的一致性查询所使用的数据块也比较小，查询的效率就会比较高。

1.1.1 高效即席查询

```
SQL> create table t
  2  (name_id,
  3  gender not null,
  4  location not null,
  5  age_group not null,
  6  data
  7  ) as
  8  select rownum,
  9  decode(ceil(dbms_random.value(0,2)),
10  1,'M',
11  2,'F')gender,
12  ceil(dbms_random.value(1,50)) location,
13  decode(ceil(dbms_random.value(0,3)),
14  1,'child',
15  2,'young',
16  3,'middle_age',
17  4,'old'),
18  rpad('*',400,'*')
19  from dual
20  connect by rownum<=100000;

Table created.

SQL> begin
  2  dbms_stats.gather_table_stats(ownname =>'SCOTT',
  3  tabname => 'T',
  4  estimate_percent => 10,
  5  method_opt=> 'for all indexed columns',
  6  cascade=>TRUE);
  7  end;
  8  /

PL/SQL procedure successfully completed.
```

```
SQL> begin
  2  dbms_stats.gather_table_stats(ownname =>'SCOTT',
  3  tabname => 'T',
  4  estimate_percent => 10,
  5  method_opt=> 'for all indexed columns',
  6  cascade=>TRUE);
  7  end;
  8  /

PL/SQL procedure successfully completed.

SQL> set lines 1000
SQL> set autotrace traceonly
SQL> select * from t where gender='M' and location in (1,10,30) and age_group='child';

647 rows selected.
```

Execution Plan

Plan hash value: 306189815

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		695	154K	685 (0)	00:00:09
1	INLIST ITERATOR					
2	TABLE ACCESS BY INDEX ROWID	T	695	154K	685 (0)	00:00:09
* 3	INDEX RANGE SCAN	IDX_UNION	695		4 (0)	00:00:01

Statistics

```
-----
      1 recursive calls
      0 db block gets
    680 consistent gets
      0 physical reads
      0 redo size
282451 bytes sent via SQL*Net to client
  857 bytes received via SQL*Net from client
   45 SQL*Net roundtrips to/from client
      0 sorts (memory)
      0 sorts (disk)
  647 rows processed
```

```
SQL> create bitmap index gender_idx on t(gender);
Index created.
SQL> create bitmap index location_idx on t(location);
Index created.
SQL> create bitmap index age_group_idx on t(age_group);
Index created.
SQL> set autot traceonly
SQL> select * from t where gender='M' and location in (1,10,30) and age_group='child';

647 rows selected.
```

Execution Plan

Plan hash value: 4134214960

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		695	154K	202 (0)	00:00:03
1	TABLE ACCESS BY INDEX ROWID	T	695	154K	202 (0)	00:00:03
2	BITMAP CONVERSION TO ROWIDS					
3	BITMAP AND					
* 4	BITMAP INDEX SINGLE VALUE	AGE_GROUP_IDX				
5	BITMAP OR					
* 6	BITMAP INDEX SINGLE VALUE	LOCATION_IDX				
* 7	BITMAP INDEX SINGLE VALUE	LOCATION_IDX				
* 8	BITMAP INDEX SINGLE VALUE	LOCATION_IDX				
* 9	BITMAP INDEX SINGLE VALUE	GENDER_IDX				

采用位图索引查看效率

Predicate Information (identified by operation id):

```
-----
4 - access("AGE_GROUP"='child')
6 - access("LOCATION"=1)
7 - access("LOCATION"=10)
8 - access("LOCATION"=30)
9 - access("GENDER"='M')
```

Statistics

```
-----
1 recursive calls
0 db block gets
637 consistent gets
40 physical reads
0 redo size
282451 bytes sent via SQL*Net to client
857 bytes received via SQL*Net from client
45 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
647 rows processed
```

1.1.2 快速统计条数

```
SQL> create table t as select * from dba_objects;
Table created.

SQL> insert into t select * from t;
9470 rows created.

SQL> /
18940 rows created.

SQL> /
37880 rows created.

SQL> /
75760 rows created.

SQL> /
151520 rows created.

SQL> /
303040 rows created.

SQL> update t set object_id=rownum;
606080 rows updated.

SQL> commit;
Commit complete.
```

```
SQL> create index idx_t_obj on t(object_id);
```

```
Index created.
```

```
SQL> alter table T modify object_id not null;
```

```
Table altered.
```

```
SQL> set autotrace traceonly
```

```
SQL> select count(*) from t;
```

```
Execution Plan
```

```
Plan hash value: 278572740
```

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	376 (3)	00:00:05
1	SORT AGGREGATE		1		
2	INDEX FAST FULL SCAN	IDX_T_OBJ	661K	376 (3)	00:00:05

同样条件下普通索引的效率

```
Statistics
```

```

205 recursive calls
0 db block gets
1473 consistent gets
1352 physical reads
0 redo size
413 bytes sent via SQL*Net to client
384 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
6 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> create bitmap index idx_bitm_t_status on t(status);
```

```
Index created.
```

```
SQL> select count(*) from t;
```

```
Execution Plan
```

```
Plan hash value: 4272013625
```

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	14 (0)	00:00:01
1	SORT AGGREGATE		1		
2	BITMAP CONVERSION COUNT		661K	14 (0)	00:00:01
3	BITMAP INDEX FAST FULL SCAN	IDX_BITM_T_STATUS			

同样条件下位图索引的查询效率


```

Statistics
-----
      5  recursive calls
      0  db block gets
    102  consistent gets
     15  physical reads
      0  redo size
    413  bytes sent via SQL*Net to client
    384  bytes received via SQL*Net from client
      2  SQL*Net roundtrips to/from client
      0  sorts (memory)
      0  sorts (disk)
      1  rows processed

```

1.2 不适合的场景

由于位图索引本身存储特性的限制，在值重复率高的列或需要经常更新的列是不适合建立位图索引的。

1.2.1 列修改容易造成锁表

```

SQL> create table test(id number,name varchar2(30),sex varchar2(10),age number);
Table created.

SQL> create bitmap index idx_test_bitmap on test(sex);
Index created.

SQL> select sid from v$mystat where rownum=1;

      SID
-----
      37

SQL> insert into test values(1,'renqinglei','man',30);
1 row created.

```

暂未提交

```

SQL> █

```

```
[oracle@std ~]$ sqlplus scott/tiger

SQL*Plus: Release 10.2.0.1.0 - Production on Thu Apr 10 10:51:11 2014

Copyright (c) 1982, 2005, Oracle. All rights reserved.

Connected to:
Oracle Database 10g Enterprise Edition Release 10.2.0.1.0 - Production
With the Partitioning, OLAP and Data Mining options

SQL> select sid from v$mystat where rownum=1;

      SID
-----
      31
```

插入等待

```
SQL> insert into test values(2,'beijing','man',20);
```

```
SQL> select sid,type,id1,id2,lmode,block from v$lock where sid in (31,37);

      SID TY          ID1          ID2          LMODE          BLOCK
-----
      37 TO           8693           1           3           0
      31 TX          1048611         55           0           0
      37 TM           17155           0           3           0
      31 TM           17155           0           3           0
      37 TX          1048611         55           6           1
      31 TX          1310731         44           6           0

6 rows selected.
```

1.2.2 列值重复度低效率低

```
SQL> create table t as select * from dba_objects;
Table created.

SQL> insert into t select * from t;
9472 rows created.

SQL> /
18944 rows created.

SQL> /
37888 rows created.

SQL> /
75776 rows created.

SQL> /
151552 rows created.

SQL> commit;
Commit complete.
```

```
SQL> create index idx_t on t(object_id);

Index created.

SQL> set autotrace traceonly
SQL> alter table t modify object_id not null;

Table altered.

SQL> select /*+index(t idx_t)*/ count(*) from t;
```

普通索引的开销

Execution Plan

Plan hash value: 3792760506

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	645 (1)	00:00:08
1	SORT AGGREGATE		1		
2	INDEX FULL SCAN	IDX_T	337K	645 (1)	00:00:08

Note

- dynamic sampling used for this statement

Statistics

```
205 recursive calls
0 db block gets
746 consistent gets
638 physical reads
```

```
SQL> set autot traceonly
SQL> select /* +index(t idx_t_bitmap)*/count(*) from t;
```

Execution Plan

Plan hash value: 435700843

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	603 (1)	00:00:08
1	SORT AGGREGATE		1		
2	BITMAP CONVERSION COUNT		1062K	603 (1)	00:00:08
3	BITMAP INDEX FAST FULL SCAN	IDX_T_BITMAP			

Note

- dynamic sampling used for this statement

Statistics

```
0 recursive calls
0 db block gets
1327 consistent gets
0 physical reads
49364 redo size
```


2. 简要说说函数索引的原理及好处

2.1 函数索引的原理

函数索引是把通过函数计算过的值存储在索引中，当利用函数进行查询时，利用计算出的函数值查询函数索引，进而查询与其函数索引值对应的 ROWID，直接通过 rowid 取出对应的记录，提高了查询效率。

2.2 使用函数索引的好处

2.2.1 对部分记录建立函数索引

```
SQL> create table t (id int ,status varchar2(2));
Table created.
SQL> create index id_normal on t(status);
Index created.
SQL> insert into t select rownum ,'Y' from dual connect by rownum<=1000000;
1000000 rows created.
SQL> commit;
Commit complete.
SQL> analyze table t compute statistics for table for all indexes for all indexed columns;
Table analyzed.
```

```
SQL> select * from t where status='N';
```

```
no rows selected
```

普通索引的开销

```
Execution Plan
```

```
Plan hash value: 2252729315
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		6667	66670	29 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	T	6667	66670	29 (0)	00:00:01
* 2	INDEX RANGE SCAN	ID_NORMAL	6667		18 (0)	00:00:01

```
Predicate Information (identified by operation id):
```

```
2 - access("STATUS"='N')
```

```
Statistics
```

```

1 recursive calls
0 db block gets
3 consistent gets
0 physical reads
0 redo size
327 bytes sent via SQL*Net to client
373 bytes received via SQL*Net from client
```

```
SQL> set autotrace off
```

```
SQL> analyze index id_normal validate structure;
```

```
Index analyzed.
```

普通索引的结构

```
SQL> select name,btree_space,lf_rows,height from index_stats;
```

NAME	BTREE_SPACE	LF_ROWS	HEIGHT
ID_NORMAL	19230668	1000000	3

```
SQL> drop index id_normal;
```

```
Index dropped.
```

```
SQL> create index id_status on t (Case when status= 'N' then 'N' end);
```

```
Index created.
```

```
SQL> analyze table t compute statistics for table for all indexes for all indexed columns;
```

```
Table analyzed.
```

```
SQL> select * from t where (case when status='N' then 'N' end)='N';
no rows selected
```

位图索引的效率

Execution Plan

Plan hash value: 1835552001

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	10	1 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	T	1	10	1 (0)	00:00:01
* 2	INDEX RANGE SCAN	ID_STATUS	1		1 (0)	00:00:01

Predicate Information (identified by operation id):

2 - access(CASE "STATUS" WHEN 'N' THEN 'N' END ='N')

Statistics

```
8 recursive calls
0 db block gets
3 consistent gets
```

```
SQL> analyze index id_status validate structure;
```

Index analyzed.

位图索引的结构

```
SQL> select name,btree_space,lf_rows,height from index_stats;
```

NAME	BTREE_SPACE	LF_ROWS	HEIGHT
ID_STATUS	7996	0	1

2.2.1 减少递归调用

```
SQL> create table t1 (first_name varchar2(200),last_name varchar2(200),id number);
Table created.

SQL> create table t2 as select * from dba_objects where rownum<=1000;
Table created.

SQL> insert into t1 (first_name,last_name,id) select object_name,object_type,rownum
  2  from dba_objects where rownum<=1000;

1000 rows created.

SQL> commit;

Commit complete.

SQL> create or replace function get_obj_name(p_id t2.object_id%type)
  2  return t2.object_name%type DETERMINISTIC is
  3  v_name t2.object_name%type;
  4  begin
  5  select object_name into v_name from t2
  6  where object_id=p_id;
  7  return v_name;
  8  end;
  9  /

Function created.
```

```
SQL> select *   from t1 where get_obj_name(id)='TEST';

no rows selected
```

没走函数索引递归调用

Statistics

```
-----
1046 recursive calls
  0 db block gets
15007 consistent gets
  0 physical reads
  0 redo size
 390 bytes sent via SQL*Net to client
 373 bytes received via SQL*Net from client
   1 SQL*Net roundtrips to/from client
   0 sorts (memory)
   0 sorts (disk)
   0 rows processed
```

```
SQL> create index idx_func_id on t1(get_obj_name(id));
Index created.
```

```
SQL> select * from t1 where get_obj_name(id)='TEST';
no rows selected
```

创建函数索引后的递归调用

```
Statistics
```

```
-----
 67 recursive calls
   0 db block gets
 16 consistent gets
   1 physical reads
   0 redo size
390 bytes sent via SQL*Net to client
373 bytes received via SQL*Net from client
   1 SQL*Net roundtrips to/from client
   0 sorts (memory)
   0 sorts (disk)
   0 rows processed
```

3.说说使用反向索引、全文索引的应用场景

3.1 反向索引的应用场景

在对一个列建立索引的时候，如果这列的值是数值型的且为连续递增型的数据，这种情况下就适合建立反向索引，因为建立普通索引将导致连续的数据在索引块中过度集中，这样虽然会带来查询上的好处但是也会造成大量的块争用，且当有多个会话修改相连的数据时，容易造成阻塞，影响系统效率。这种情况下建立反向索引后，根据反向索引的特性，在插入索引块之前，系统会把索引列值倒过来插入，从而避免连续的数据插入相同的数据块中，如 123、124、125、126 倒过来以后就会变为 321、421、521、621。反向索引的一大缺点是无法做范围查询，试想倒过来的值已经无法判断他们之间真是的范围了。

3.2 全文索引的应用场景

全文索引适合于在一些大字段类型中查找匹配关键字，例如搜索引擎（谷歌、百度）常会用到。

3.2.1 全文索引环境准备

```
SQL> select role from dba_roles where role='CTXAPP';
no rows selected

SQL> select username from dba_users where username='CTXSYS';
no rows selected
```

```
SQL> create tablespace idx_ctxsys
  2  datafile '/u02/app/oradata/PSDB/disk1/idx_ctxsys01.dbf' size 200m;
Tablespace created.

SQL> @?/ctx/admin/catctx.sql ctxsys idx_ctxsys TEMPTS1 nolog
```

```
SQL> select username,account_status from dba_users where username='CTXSYS';

USERNAME                                ACCOUNT_STATUS
-----
CTXSYS                                  OPEN

SQL> select role from dba_roles where role='CTXAPP';

ROLE
-----
CTXAPP
```

执行脚本后查询用户和角色已创建

```
SQL> conn ctxsys/ctxsys
Connected.
SQL> @?/ctx/admin/defaults/drdefus.sql
Creating lexer preference...

PL/SQL procedure successfully completed.

Creating wordlist preference...

PL/SQL procedure successfully completed.

Creating stoplist...

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

Creating default policy...

PL/SQL procedure successfully completed.

SQL>
```

3.2.2 创建语法分析器并设置属性

```
SQL> conn ctxsys/ctxsys
Connected.
SQL> grant execute on ctx_ddl to scott;

Grant succeeded.

SQL> conn scott/tiger
Connected.
SQL> ctx_ddl.create_preference('club_lexer','CHINESE_LEXER');
SP2-0734: unknown command beginning "ctx_ddl.cr..." - rest of line ignored.
SQL> exec ctx_ddl.create_preference('club_lexer','CHINESE_LEXER');

PL/SQL procedure successfully completed.

SQL> exec ctx_ddl.create_preference('mywordlist', 'BASIC_WORDLIST');

PL/SQL procedure successfully completed.

SQL> exec ctx_ddl.set_attribute('mywordlist','PREFIX_INDEX','TRUE');

PL/SQL procedure successfully completed.

SQL> exec ctx_ddl.set_attribute('mywordlist','PREFIX_MIN_LENGTH',1);

PL/SQL procedure successfully completed.

SQL> exec ctx_ddl.set_attribute('mywordlist','PREFIX_MAX_LENGTH', 5);

PL/SQL procedure successfully completed.

SQL> exec ctx_ddl.set_attribute('mywordlist','SUBSTRING_INDEX', 'YES');

PL/SQL procedure successfully completed.
```

3.2.3 创建表及为表创建全文索引

```
SQL> create table test as select * from dba_objects;

Table created.

SQL> update test set object_name='高兴' where rownum<=2;

2 rows updated.

SQL> commit;

Commit complete.

SQL> create index id_cont_test on TEST (object_name) indextype is ctxsys.context
  2  parameters ('DATASTORE CTXSYS.DIRECT DATASTORE FILTER
  3  CTXSYS.NULL_FILTER LEXER club_lexer WORDLIST mywordlist');

Index created.

SQL> exec ctx_ddl.sync_index('id_cont_TEST', '20M');

PL/SQL procedure successfully completed.
```

3.2.4 比较全文索引的性能

```
SQL> set linesize 1000
SQL> set autotrace traceonly
SQL> select * from test where OBJECT_NAME like '%高兴%';
```

Execution Plan

Plan hash value: 1357081020

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		4	708	36 (0)	00:00:01
* 1	TABLE ACCESS FULL	TEST	4	708	36 (0)	00:00:01

Predicate Information (identified by operation id):

```
1 - filter("OBJECT_NAME" LIKE '%????%')
```

Note

```
- dynamic sampling used for this statement
```

Statistics

```
436 recursive calls
0 db block gets
251 consistent gets
0 physical reads
```

SQL> select * from test where contains(OBJECT_NAME, '高兴')>0;

Execution Plan

Plan hash value: 1428524535

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		5	945	5 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	TEST	5	945	5 (0)	00:00:01
* 2	DOMAIN INDEX	ID_CONT_TEST			4 (0)	00:00:01

Predicate Information (identified by operation id):

```
2 - access("CTXSYS"."CONTAINS"("OBJECT_NAME", '????')>0)
```

Note

```
- dynamic sampling used for this statement
```

Statistics

```
2859 recursive calls
0 db block gets
2200 consistent gets
1 physical reads
```

我们知道像这种方式是无论如何都用不上索引的

适用全文索引的写法

用到了全文索引，cost值才5，一致性读有点大，但系统时通过cost来决定性能的